



Deep Learning


MPHY0041 Machine Learning in Medical Imaging

Yipeng Hu
yipeng.hu@ucl.ac.uk

Network Architecture | Width and Depth

Network Architecture | Width and Depth





Output layer - l

Input \rightarrow H1: $i \rightarrow j$

$$z_j^1 = g_j^1 \left[(\mathbf{w}_j^0)^T \mathbf{x} + b_j^1 \right]$$

H1 \rightarrow H2: $j \rightarrow k$

$$z_k^2 = g_k^1 \left[(\mathbf{w}_k^1)^T \mathbf{z}^1 + b_k^1 \right]$$

H2 \rightarrow Output: $k \rightarrow l$

$$y_k = g_l^1 \left[(\mathbf{w}_l^2)^T \mathbf{z}^2 + b_l^2 \right]$$

Universal Approximation Theorem



Any continuous functions on bounded and closed functions or between finite dimensional discrete space

Width and Depth

Efficiency

Hierarchical representation learning

Empirical results




An example using an absolute value rectification

Network Architecture | Branching, Joining and Skipping

Branching and Joining

“Inception”



Multi-stream

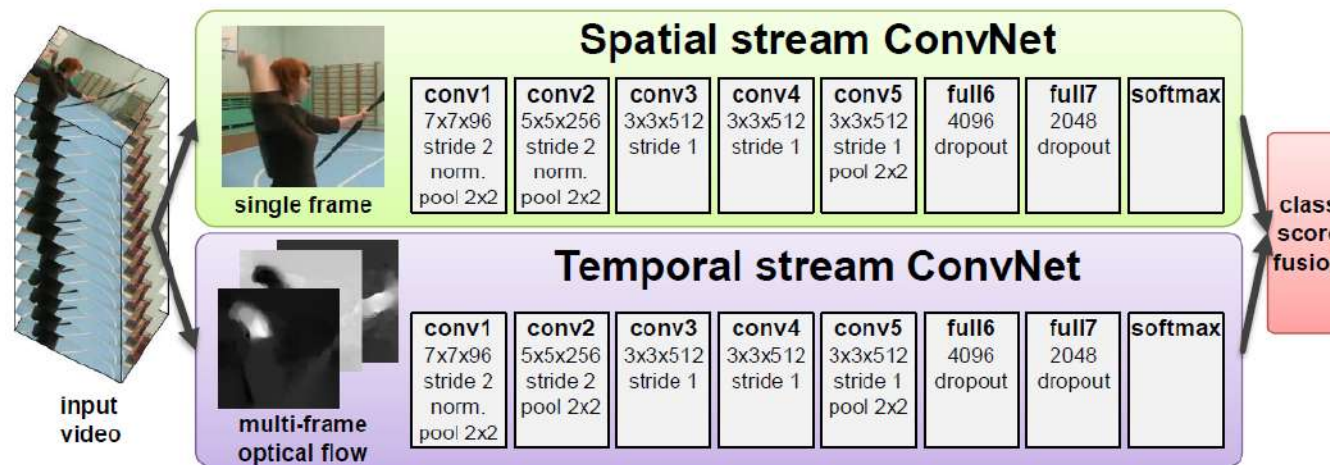
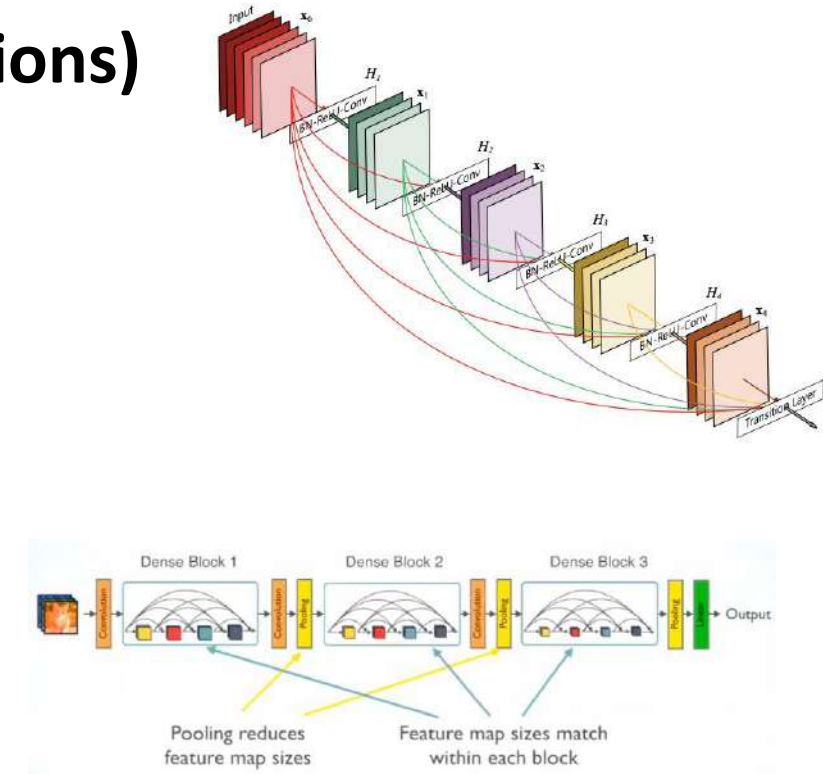
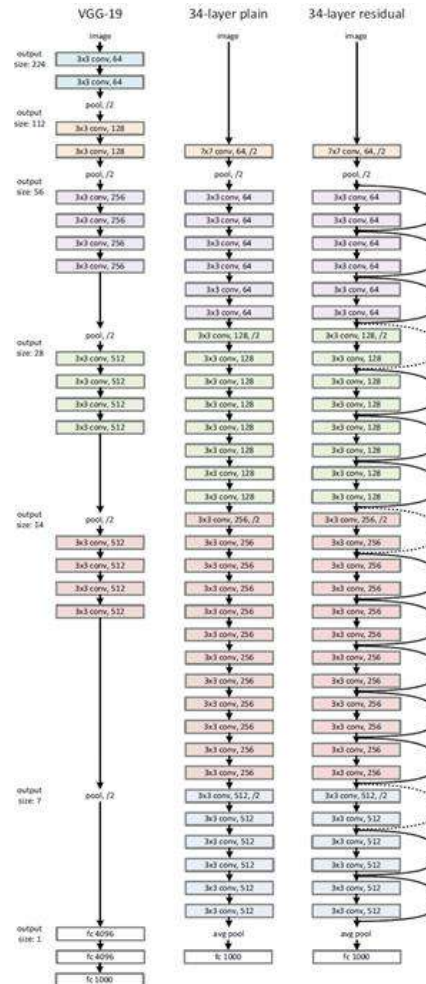
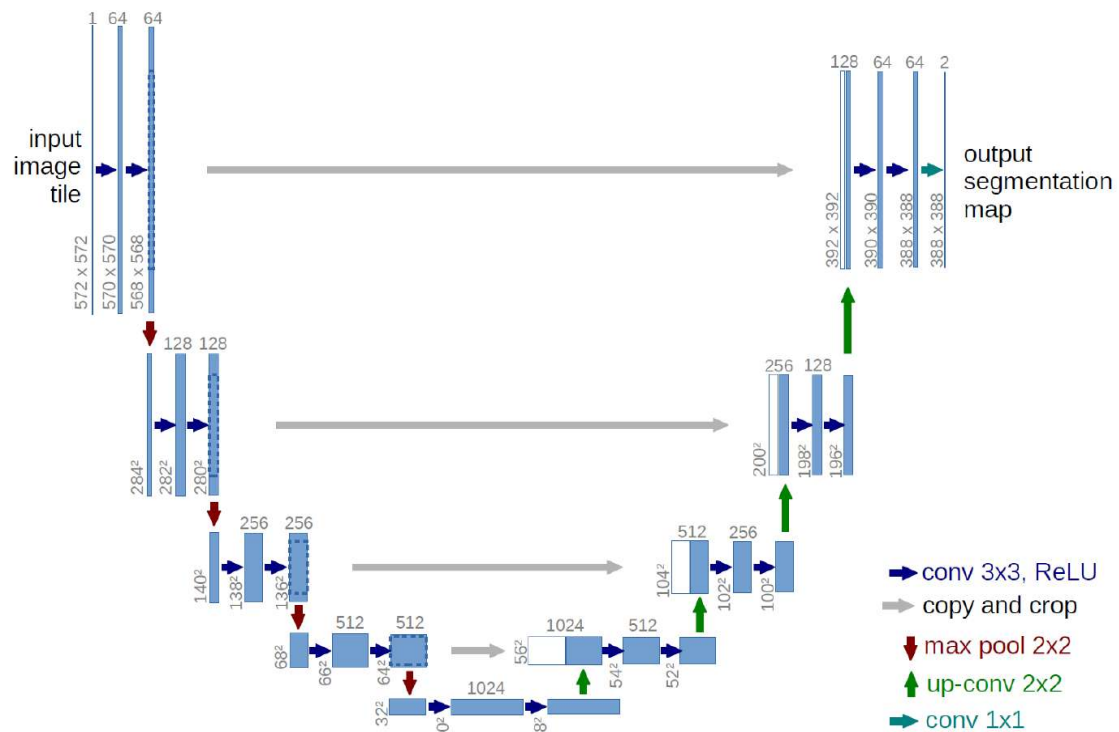



Figure 1: Two-stream architecture for video classification.

Multi-task learning*


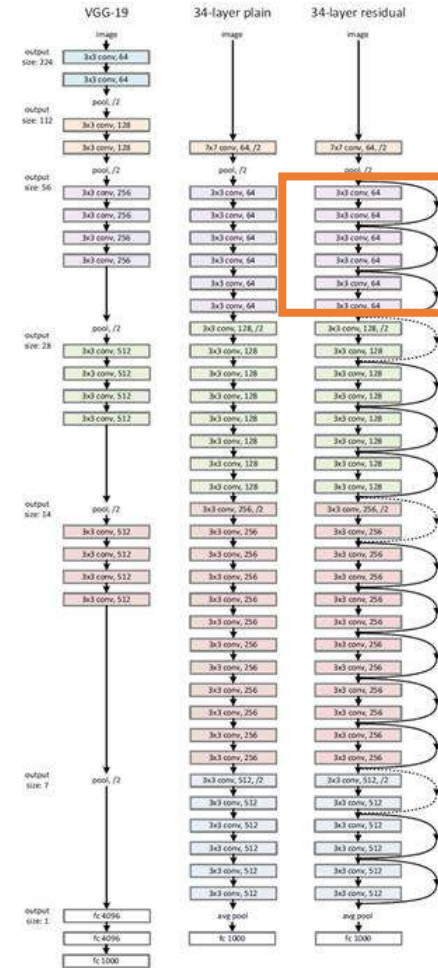
Skipping (shortcuts, skip layers and residual connections)




The ResNet



$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x}.$$



Network Architecture | Convolutional Neural Networks



Output layer - l

Input \rightarrow H1: $i \rightarrow j$

$$z_j^1 = g_j^1 \left[(\mathbf{w}_j^0)^T \mathbf{x} + b_j^1 \right]$$

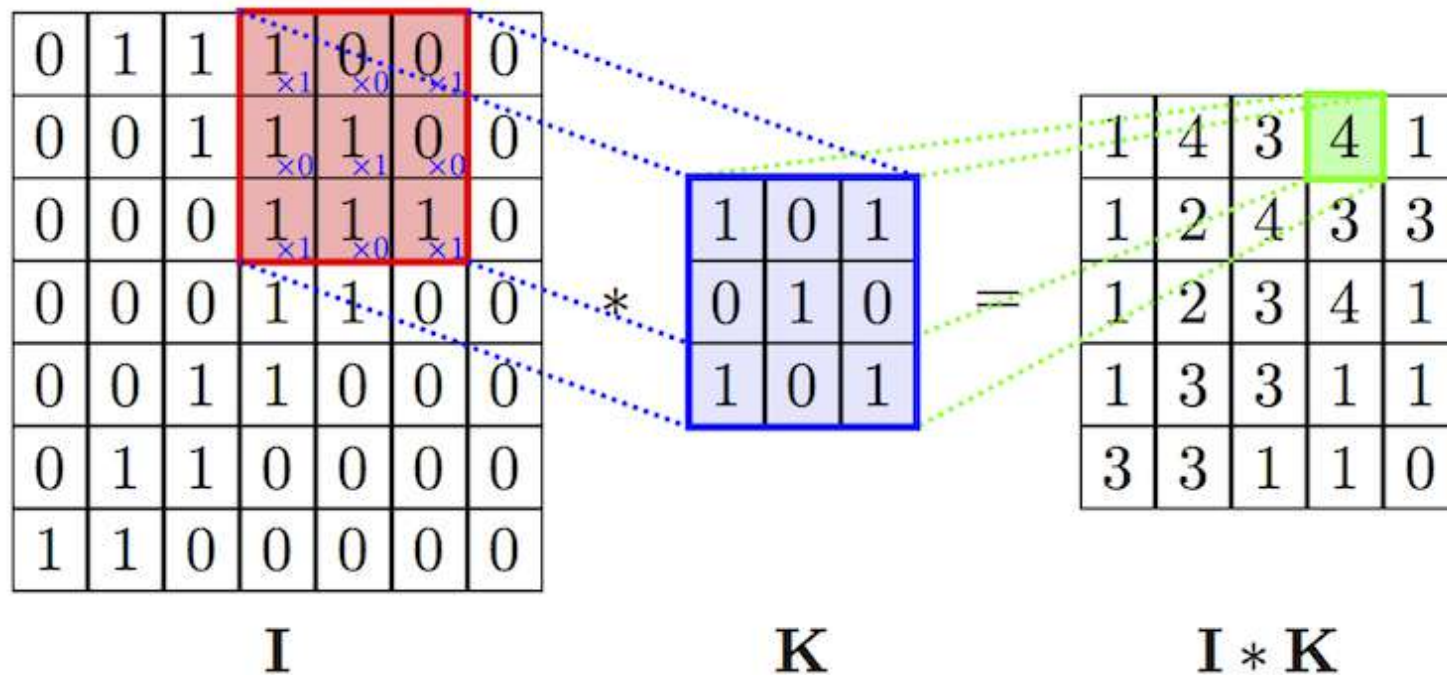
H1 \rightarrow H2: $j \rightarrow k$

$$z_k^2 = g_k^1 \left[(\mathbf{w}_k^1)^T \mathbf{z}^1 + b_k^1 \right]$$


H2 \rightarrow Output: $k \rightarrow l$

$$y_k = g_l^1 \left[(\mathbf{w}_l^2)^T \mathbf{z}^2 + b_l^2 \right]$$

Convolutional Neural Networks



Convolutional Layer



Convolutional Neural Networks

$$\begin{pmatrix} x1 & x2 & x3 \\ x4 & x5 & x6 \\ x7 & x8 & x9 \end{pmatrix} \begin{pmatrix} k1 & k2 \\ k3 & k4 \end{pmatrix} = \begin{pmatrix} k1x1 + k2x2 + k3x4 + k4x5 \\ k1x2 + k2x3 + k3x5 + k4x6 \\ k1x4 + k2x5 + k3x7 + k4x8 \\ k1x5 + k2x6 + k3x8 + k4x9 \end{pmatrix}$$

$$= \begin{pmatrix} k1 & k2 & 0 & k3 & k4 & 0 & 0 & 0 & 0 \\ 0 & k1 & k2 & 0 & k3 & k4 & 0 & 0 & 0 \\ 0 & 0 & 0 & k1 & k2 & 0 & k3 & k4 & 0 \\ 0 & 0 & 0 & 0 & k1 & k2 & 0 & k3 & k4 \end{pmatrix} \cdot \begin{pmatrix} x1 \\ x2 \\ x3 \\ x4 \\ x5 \\ x6 \\ x7 \\ x8 \\ x9 \end{pmatrix}$$

Convolutional Layer

Input $\rightarrow H1: i \rightarrow j$

$$z_j^1 = \sigma_j^1 \left[(\mathbf{w}_j^0)^T \mathbf{x} + b_j^1 \right]$$


H1 $\rightarrow H2: j \rightarrow k$

$$z_k^2 = \sigma_k^1 \left[(\mathbf{w}_k^1)^T \mathbf{z}^1 + b_k^1 \right]$$



H2 \rightarrow Output: $k \rightarrow l$


$$y_k = \sigma_l^1 \left[(\mathbf{w}_l^2)^T \mathbf{z}^2 + b_l^2 \right]$$

Convolutional Layers + Fully Connected Layers



Convolutional Layers + Fully Connected Layers





Network Architecture | Sampling

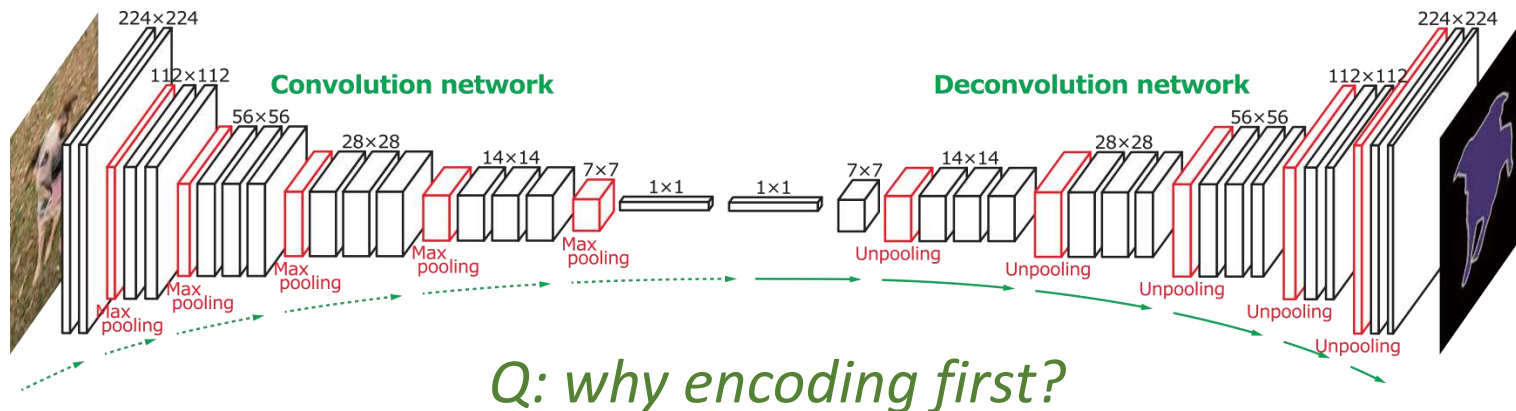

Down-sampling and up-sampling (Encoding and Decoding)

AlexNet / VGG




Fully convolutional

Global and local pooling

Dense prediction




Down-sampling and up-sampling 2D/3D feature maps



Convolution and pooling with strides

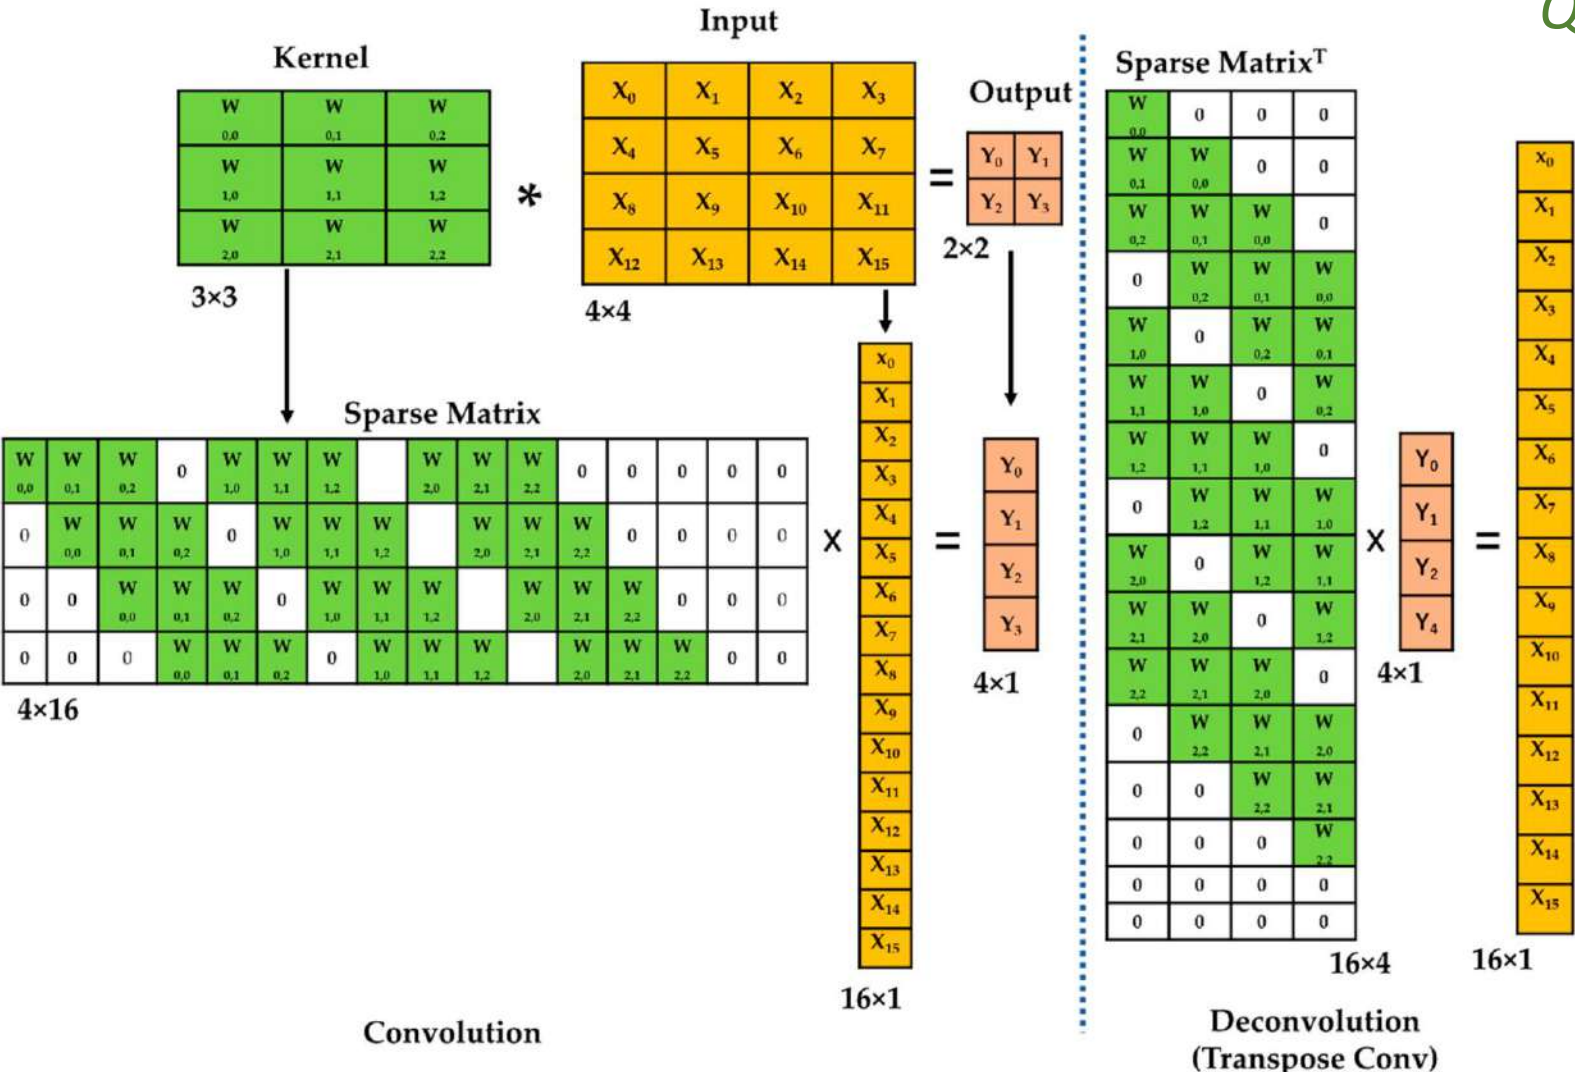
Transpose convolution and “un-pooling”

Linear sampling




Transpose convolution

Q: is transpose convolution necessary?




Network Architecture | Examples

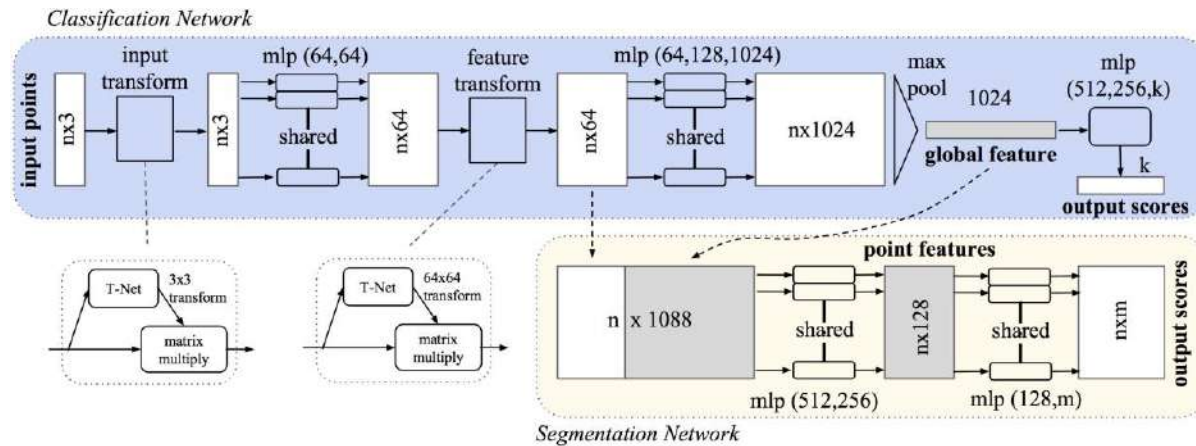
Autoencoder



UNet

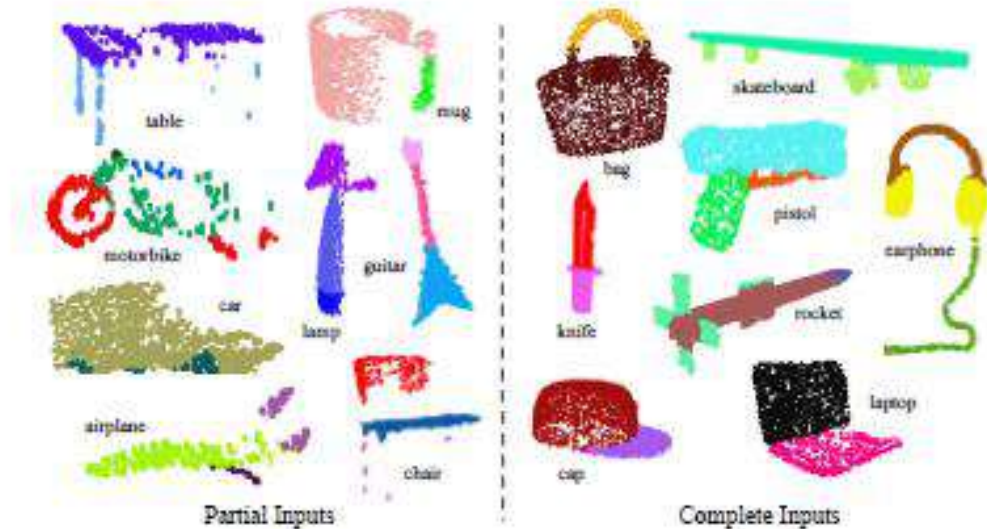



PointNet

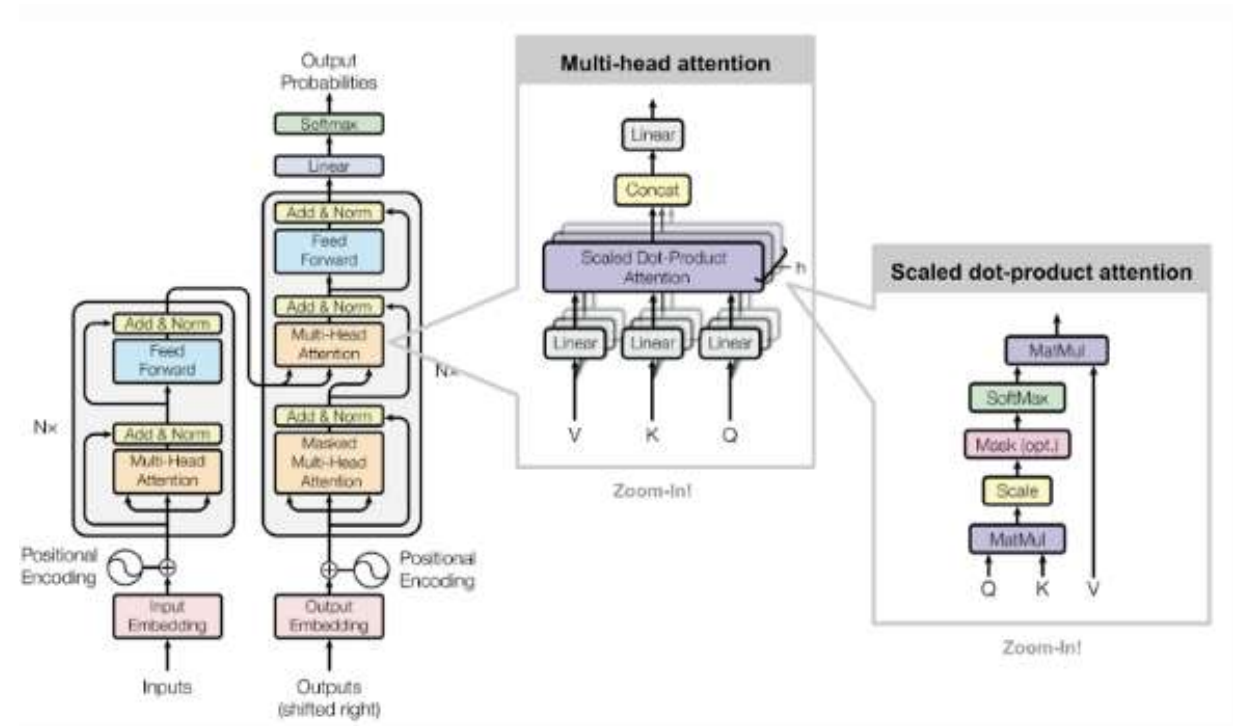
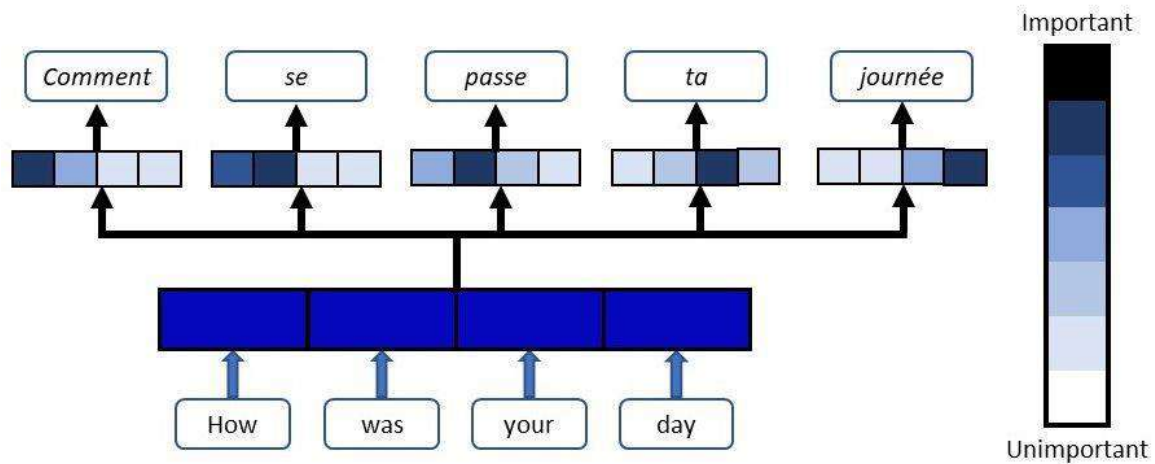


$$f(\{x_1, \dots, x_n\}) \approx g(h(x_1), \dots, h(x_n))$$

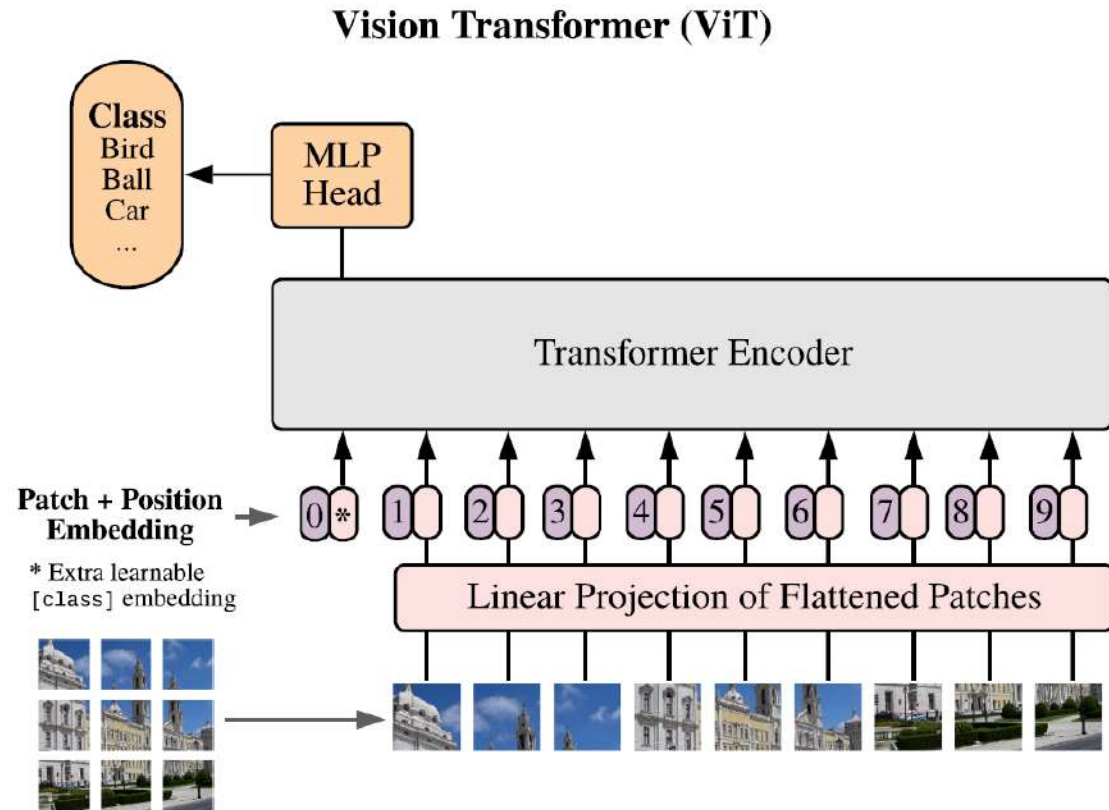
Permutation-Invariant Functions




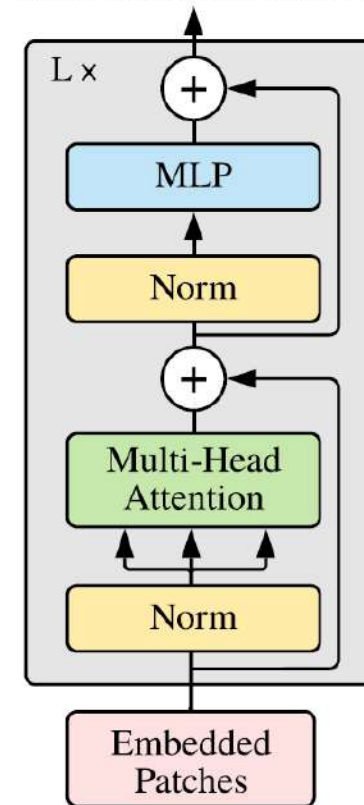
Attention Mechanism



Vision Transformer






Transformer Encoder



Network Architecture | Recurrent Neural Networks


Sequential data



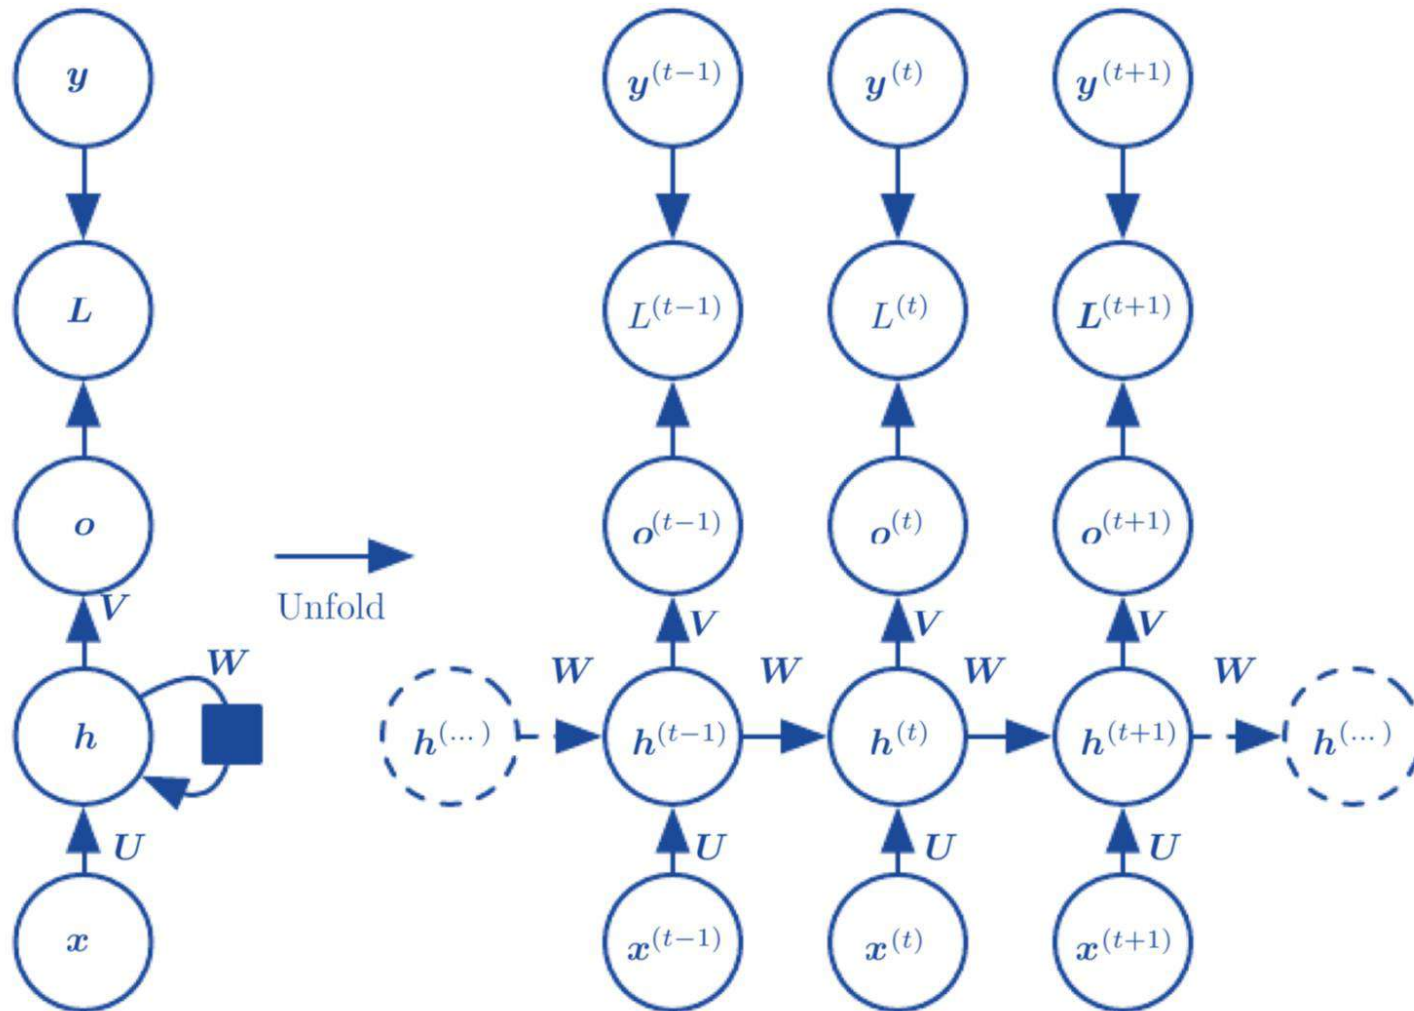
$$p(\mathbf{x}_1, \dots, \mathbf{x}_N) = \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$$

Q: examples in medical imaging?

Unfolding neural networks




Unfolding neural networks




(Truncated) backpropagation through time (BPTT)


one to one




one to many




many to one




many to many




many to many



Long short-term memory (LSTM)



LSTM (gated methods and cell state, “a belt for information flow”)



Standard RNN

How to design and choose

Application requirements

- Input / output dimensionalities
- Data format, feature vectors, images, time series

Practice

- Baseline and ablation experiments
- “Do no harm” principles
- Complicated relationship with data and training strategies
- An art?